

CS 251/340 - Machine Learning

Spring 2019, AUA

Homework No. 03

Program files submission - Due date: 23:55, March 13, 2019

1. Use the `rnorm()` function to generate a predictor X of length $n = 100$, as well as a noise vector e of length $n = 100$ (**score = 5**). Generate a response vector Y of length $n = 100$ according to the model

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + e,$$

where $\beta_0, \beta_1, \beta_2, \beta_3$ are constants of your choice (**score = 5**).

2. Fit the **Ridge** model to the simulated data, using X, X^2, \dots, X^{10} as predictors (**score = 10**). Use cross-validation (10-fold) to select the optimal value of λ (**score = 10**). Create plots of the cross-validation error as a function of λ (**score = 10**). Report the resulting coefficient estimates and discuss the results (**score = 10**).
3. Fit the **Lasso** model to the simulated data, using X, X^2, \dots, X^{10} as predictors (**score = 10**). Use cross-validation (10-fold) to select the optimal value of λ (**score = 10**). Create plots of the cross-validation error as a function of λ (**score = 10**). Report the resulting coefficient estimates and discuss the results (**score = 10**).
4. Compare the results (**score = 10**).